

Alex Vance - Resume

CAREER OBJECTIVE

To lead innovative AI engineering teams in developing high-throughput LLM architectures and production RAG retrieval platforms.

EXECUTIVE SUMMARY

Driven Senior AI Engineer specializing in LLM production systems, RAG retrieval architectures, and high-concurrency async Python services with a track record of scaling high-throughput applications.

EXPERIENCE

Lead AI Engineer | NeuralTech AI (2022 - Present)

Architected enterprise RAG document retrieval engines using FAISS & Groq API, scaling query throughput by 300%.

Engineered async FastAPI microservices handling 50k+ daily streaming requests.

Software Engineer | DataFlow Systems (2019 - 2022)

Developed React & TypeScript dashboards for real-time model monitoring.

Optimized SQL queries reducing analytics latency by 45%.

EDUCATION

B.S. in Computer Science - University of California, Berkeley (2019)

TECHNICAL SKILLS

Python, FastAPI, Groq API, LangChain, PyTorch, React, TypeScript, Tailwind CSS, PostgreSQL, Docker, Git, Vector Databases, RAG Pipeline Optimization, Prompt Engineering, Microservices Design

PROJECTS

NovaChat AI Platform: Production AI Platform with real-time SSE token streaming, multi-format file RAG Q&A, and single-page PDF generator.

CERTIFICATIONS & ACHIEVEMENTS

AWS Certified Machine Learning Specialist, TensorFlow Developer Certificate

Winner of Global AI Innovation Hackathon (1st place out of 400 teams)

Published research paper on Context Window Compression in LLMs

LANGUAGES

English (Native), Spanish (Professional), Japanese (Conversational)